



CS 145 Midterm Review

Materials to know

| | |
|------------------------------|---|
| Know your SQL | Select-Project-Aggregate-Having, Nested queries, NULL values Set vs multi-sets |
| How to design 'good' tables? | Find bad FDs, Compute closures/superkeys, Do BCNF decompositions, Find MVDs |
| ACID transactions | Performance #s (disk seeks, memory access) WAL logs for a transaction? How to recover post crash? (Recap: WAL & performance, Why is a log COMMIT fast?) |
| | Is a schedule conflict serializable? How to build Conflict graph? What happens during 2PL execution? (Recap: With a 5 Transaction example) |

1. Summary recap on a few key topics (see main lectures !!)
2. Mainly, honing a few popular questions in piazza/OH

(Start with transactions and go backwards)

Transactions

Summary

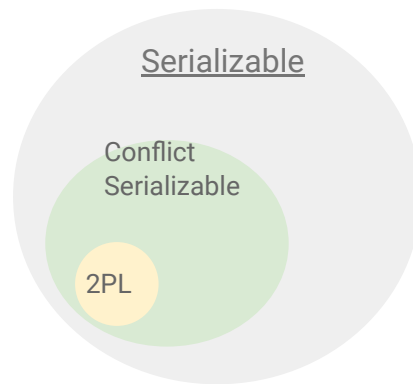
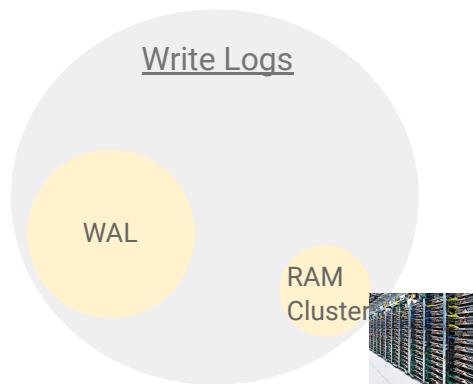
Why study Transactions?

Good programming model for parallel applications on shared data !

Atomic
Consistent
Isolation
Durable

Design choices?

- Write update Logs (e.g., WAL logs)
- Serial? Parallel, interleaved and serializable?



"Need to Master Extreme Transactions"
([Forbes \(Insights\)](#))



Latency numbers every
engineer should know

Ballpark timings

| | |
|-------------------------------------|--|
| execute typical instruction | 1/1,000,000,000 sec = 1 nanosec |
| fetch from L1 cache memory | 0.5 nanosec |
| fetch from L2 cache memory | 7 nanosec |
| Mutex lock/unlock | 25 nanosec |
| fetch from main memory | 100 nanosec |
| send 2K bytes over 1Gbps network | 20,000 nanosec |
| read 1MB sequentially from memory | 250,000 nanosec |
| fetch from new disk location (seek) | 8,000,000 nanosec |
| read 1MB sequentially from disk | 20,000,000 nanosec |
| send packet US to Europe and back | 150 milliseconds = 150,000,000 nanosec |

(~0.25 msecs)

(~10 msecs)

(~20 msecs)

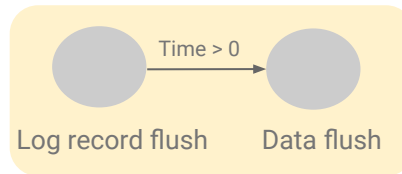
Write-Ahead Logging (WAL)

Algorithm: WAL

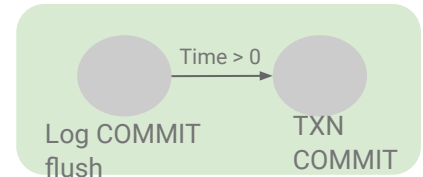
1. Must *force log record* for an update *before* the corresponding data page goes to storage
2. Must write all log records for a TX before commit

→ **Atomicity**

→ **Durability**



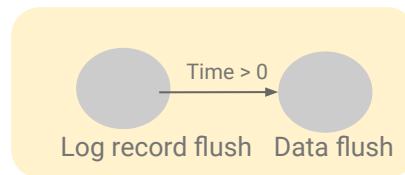
For each record update



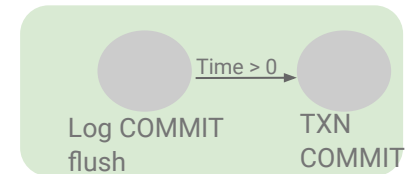
At end of TXN

Example WAL scenarios

| | |
|---|--|
| TXN commit before Log COMMIT on disk? | No |
| Data page on disk before Log flush for that record? | No |
| TXN commit before data page flushed to disk? | Yes, often. Especially for large transactions. For TXN Commit, should have Flushed... <ul style="list-style-type: none">- All record updates to Log- COMMIT record to Log |
| Later transaction needing updated data, not flushed to disk | Check memory for latest. If not in memory, read latest from disk |
| Parallel transaction looking at updated data, not flushed to disk | (Later, 2PL or Locking kicks in for RW, WW, WR conflicts) |



For each record update



At end of TXN

Example

Monthly bank interest transaction

Money

| Account | | Balance (\$) |
|---------|------|--------------|
| 3001 | | 500 |
| 4001 | | 100 |
| 5001 | | 20 |
| 6001 | | 60 |
| 3002 | | 80 |
| 4002 | | -200 |
| 5002 | | 320 |
| ... | | ... |
| 30108 | | -100 |
| 40008 | | 100 |
| 50002 | | 20 |

Money (@4:29 am day+1)

| Account | | Balance (\$) |
|---------|------|--------------|
| 3001 | | 550 |
| 4001 | | 110 |
| 5001 | | 22 |
| 6001 | | 66 |
| 3002 | | 88 |
| 4002 | | -220 |
| 5002 | | 352 |
| ... | | ... |
| 30108 | | -110 |
| 40008 | | 110 |
| 50002 | | 22 |

WAL (@4:29 am day+1)

| | | | |
|---------------|--------------------------|------|------|
| T-Monthly-423 | START TRANSACTION | | |
| T-Monthly-423 | 3001 | 500 | 550 |
| T-Monthly-423 | 4001 | 100 | 110 |
| T-Monthly-423 | 5001 | 20 | 22 |
| T-Monthly-423 | 6001 | 60 | 66 |
| T-Monthly-423 | 3002 | 80 | 88 |
| T-Monthly-423 | 4002 | -200 | -220 |
| T-Monthly-423 | 5002 | 320 | 352 |
| T-Monthly-423 | ... | ... | ... |
| T-Monthly-423 | 30108 | -100 | -110 |
| T-Monthly-423 | 40008 | 100 | 110 |
| T-Monthly-423 | 50002 | 20 | 22 |
| T-Monthly-423 | COMMIT | | |

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

START TRANSACTION

UPDATE Money

SET Balance = Balance * 1.1

COMMIT

Example

Monthly
bank
interest
transaction

With crash

Money

| Account | | Balance (\$) |
|---------|------|--------------|
| 3001 | | 500 |
| 4001 | | 100 |
| 5001 | | 20 |
| 6001 | | 60 |
| 3002 | | 80 |
| 4002 | | -200 |
| 5002 | | 320 |
| ... | | ... |
| 30108 | | -100 |
| 40008 | | 100 |
| 50002 | | 20 |

Money (@10:45 am)

| Account | | Balance (\$) |
|---------|------|--------------|
| 3001 | | 550 |
| 4001 | | 110 |
| 5001 | | 22 |
| 6001 | | 66 |
| 3002 | | 88 |
| 4002 | | -200 |
| 5002 | | 320 |
| ... | | ... |
| 30108 | | -110 |
| 40008 | | 110 |
| 50002 | | 22 |

WAL log (@10:29 am)

| T-Monthly-423 | START TRANSACTION | | |
|---------------|-------------------|------|------|
| T-Monthly-423 | 3001 | 500 | 550 |
| T-Monthly-423 | 4001 | 100 | 110 |
| T-Monthly-423 | 5001 | 20 | 22 |
| T-Monthly-423 | 6001 | 60 | 66 |
| T-Monthly-423 | 3002 | 80 | 88 |
| T-Monthly-423 | ... | ... | ... |
| T-Monthly-423 | 30108 | -100 | -110 |
| T-Monthly-423 | 40008 | 100 | 110 |
| T-Monthly-423 | 50002 | 20 | 22 |
| T-Monthly-423 | 4002 | -200 | -220 |
| T-Monthly-423 | 5002 | 320 | 352 |

??

??

??

??

'T-Monthly-423'

Monthly Interest 10%

4:28 am Starts run on 10M bank accounts

Takes 24 hours to run

Network outage at 10:29 am,

System access at 10:45 am

Did T-Monthly-423 complete?

Which tuples are bad?

Case1: T-Monthly-423 was crashed

Case2: T-Monthly-423 completed. 4002

deposited 20\$ at 10:45 am

Example

Monthly bank interest transaction

Recovery

Money (@10:45 am)

| Account | ... | Balance (\$) |
|---------|-----|--------------|
| 3001 | | 550 |
| 4001 | | 110 |
| 5001 | | 22 |
| 6001 | | 66 |
| 3002 | | 88 |
| 4002 | | -200 |
| 5002 | | 320 |
| ... | | |
| 30108 | | -110 |
| 40008 | | 110 |
| 50002 | | 22 |

Money (after recovery)

| Account | ... | Balance (\$) |
|---------|-----|--------------|
| 3001 | | 500 |
| 4001 | | 100 |
| 5001 | | 20 |
| 6001 | | 60 |
| 3002 | | 80 |
| 4002 | | -200 |
| 5002 | | 320 |
| ... | | ... |
| 30108 | | -100 |
| 40008 | | 100 |
| 50002 | | 20 |

WAL log (@10:29 am)

| | | | |
|---------------|-------------------|------|------|
| T-Monthly-423 | START TRANSACTION | | |
| T-Monthly-423 | 3001 | 500 | 550 |
| T-Monthly-423 | 4001 | 100 | 110 |
| T-Monthly-423 | 5001 | 20 | 22 |
| T-Monthly-423 | 6001 | 60 | 66 |
| T-Monthly-423 | 3002 | 80 | 88 |
| T-Monthly-423 | ... | ... | ... |
| T-Monthly-423 | 30108 | -100 | -110 |
| T-Monthly-423 | 40008 | 100 | 110 |
| T-Monthly-423 | 50002 | 20 | 22 |
| T-Monthly-423 | 4002 | -200 | -220 |
| T-Monthly-423 | 5002 | 320 | 352 |

System recovery (after 10:45 am)

- 1 Rollback uncommitted transactions
 - Restore old values from WALlog (if any)
 - Notify developers about aborted txn
- 1.1 Redo Recent committed transactions (w/ new values)
- 2 Back in business
- 3 Redo (any pending) transactions

(Sometimes swap 2 and 3, as a tradeoff)

Example

Monthly bank interest transaction

Performance

Money

| Account | | Balance (\$) |
|---------|------|--------------|
| 3001 | | 500 |
| 4001 | | 100 |
| 5001 | | 20 |
| 6001 | | 60 |
| 3002 | | 80 |
| 4002 | | -200 |
| 5002 | | 320 |
| ... | | ... |
| 30108 | | -100 |
| 40008 | | 100 |
| 50002 | | 20 |

Money (@4:29 am day+1)

| Account | | Balance (\$) |
|---------|------|--------------|
| 3001 | | 550 |
| 4001 | | 110 |
| 5001 | | 22 |
| 6001 | | 66 |
| 3002 | | 88 |
| 4002 | | -220 |
| 5002 | | 352 |
| ... | | ... |
| 30108 | | -110 |
| 40008 | | 110 |
| 50002 | | 22 |

WAL (@4:29 am day+1)

| | | | |
|---------------|--------------------------|------|------|
| T-Monthly-423 | START TRANSACTION | | |
| T-Monthly-423 | 3001 | 500 | 550 |
| T-Monthly-423 | 4001 | 100 | 110 |
| T-Monthly-423 | 5001 | 20 | 22 |
| T-Monthly-423 | 6001 | 60 | 66 |
| T-Monthly-423 | 3002 | 80 | 88 |
| T-Monthly-423 | 4002 | -200 | -220 |
| T-Monthly-423 | 5002 | 320 | 352 |
| T-Monthly-423 | ... | ... | ... |
| T-Monthly-423 | 30108 | -100 | -110 |
| T-Monthly-423 | 40008 | 100 | 110 |
| T-Monthly-423 | 50002 | 20 | 22 |
| T-Monthly-423 | COMMIT | | |

Cost to update all data

10M bank accounts → 10M seeks? (worst case)

(@10 msec/seek, that's 100,000 secs)



Cost to Append to log

- + 1 seek to get 'end of log'
- + write 10M log entries sequentially (fast!)
(< 1 sec)

[Lazily update data on disk later, when convenient.]

Materials to
know

A few
questions

| | |
|------------------------------|---|
| Know your SQL | Select-Project-Aggregate-Having, Nested queries, NULL values Set vs multi-sets |
| How to design 'good' tables? | Find bad FDs, Compute closures/superkeys, Do BCNF decompositions, Find MVDs |
| ACID transactions | Performance #s (disk seeks, memory access) WAL logs for a transaction? How to recover post crash? (Recap: WAL & performance, Why is a log COMMIT fast?) |
| | Is a schedule conflict serializable? How to build Conflict graph? What happens during 2PL execution? (Recap: With a 5 Transaction example) |

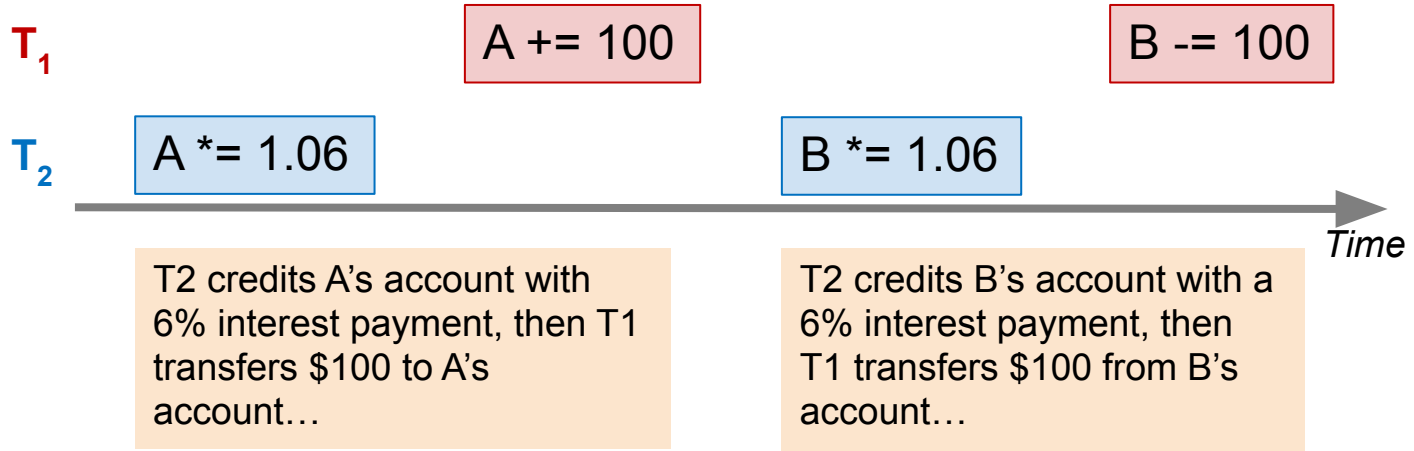
(Recap will start with transactions and go backwards)

Example- consider two TXNs:

T_1 $A += 100$ $B -= 100$

T_2 $A *= 1.06$ $B *= 1.06$

The DBMS can **interleave** the TXNs





Scheduling Definitions

- A **serial schedule** is one that does not interleave the actions of different transactions
- A and B are **equivalent schedules** if, ***for any database state***, the effect on DB of executing A **is identical to** the effect of executing B
- A **serializable schedule** is a schedule that is equivalent to ***some*** serial execution of the transactions.

The word “**some**” makes this def powerful and tricky!



Conflict Types

Two actions **conflict** if they are part of different TXNs, involve the same variable, and at least one of them is a write

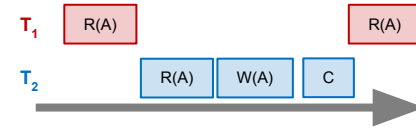
- Thus, there are three types of conflicts:
 - Read-Write conflicts (RW)
 - Write-Read conflicts (WR)
 - Write-Write conflicts (WW)

Why no “RR Conflict”?

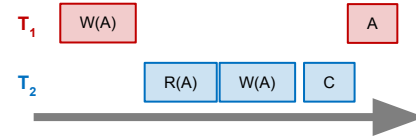
Interleaving anomalies occur with / because of these conflicts between TXNs (*but these conflicts can occur without causing anomalies!*)

Classic Anomalies with Interleaved Execution

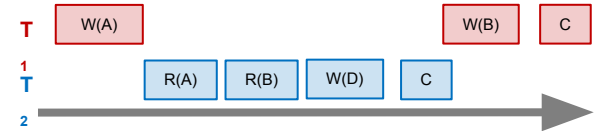
“Unrepeatable read”:



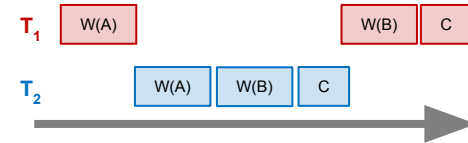
“Dirty read” / Reading uncommitted data:



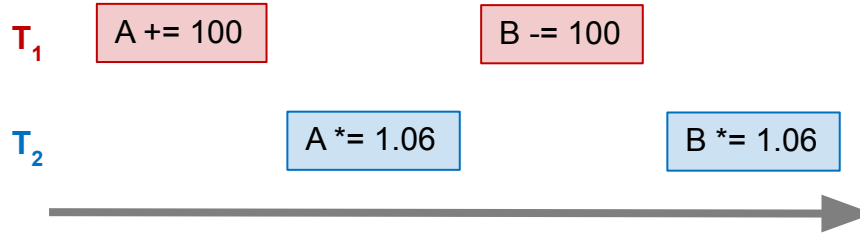
“Inconsistent read” / Reading partial commits:



Partially-lost update:



Serializable?



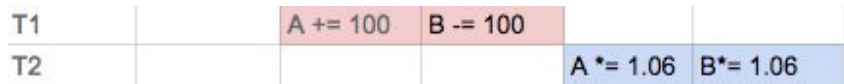
Serial schedules:

| | A | B |
|-----------------------|--------------------|--------------------|
| $T_1 \rightarrow T_2$ | $1.06 * (A + 100)$ | $1.06 * (B - 100)$ |
| $T_2 \rightarrow T_1$ | $1.06 * A + 100$ | $1.06 * B - 100$ |

| A | B |
|--------------------|--------------------|
| $1.06 * (A + 100)$ | $1.06 * (B - 100)$ |

Same as a serial schedule
for all possible values of
 $A, B = \text{serializable}$

Serial Schedules

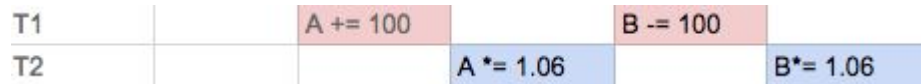


S1

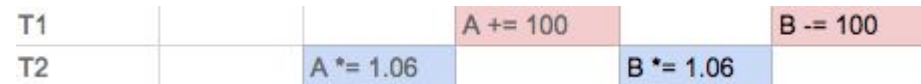


S2

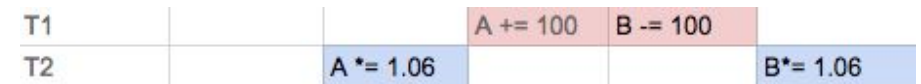
Interleaved Schedules



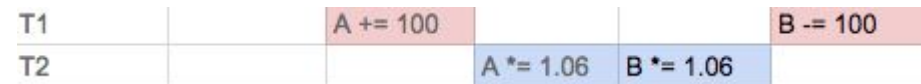
S3



S4



S5



S6

| | |
|----------------------------------|----------------------|
| Serial Schedules | S1, S2 |
| Serializable Schedules | S3, S4 |
| Equivalent Schedules | <S1, S3> <S2, S4> |
| Non-serializable (Bad) Schedules | S5, S6 |

Example with 5 Transactions

Schedule S1

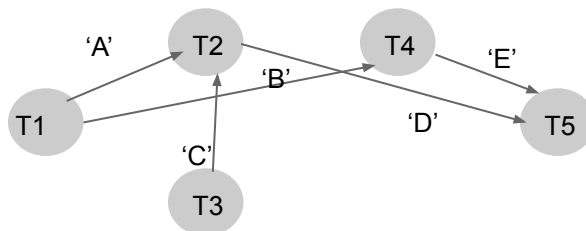
| | | | | | | | | | |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| w1(A) | r2(A) | w1(B) | w3(C) | r2(C) | r4(B) | w2(D) | w4(E) | r5(D) | w5(E) |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|

Good or Bad schedule?
Conflict serializable?

Step1

Find conflicts
(RW, WW, WR)

| | | | | | | | | | |
|----|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| T1 | w1(A) | | w1(B) | | | | | | |
| T2 | | r2(A) | | r2(C) | | w2(D) | | | |
| T3 | | | | w3(C) | | | | | |
| T4 | | | | | r4(B) | | w4(E) | | |
| T5 | | | | | | | | r5(D) | w5(E) |



Acyclic
⇒ Conflict serializable!
⇒ Serializable

Step2

Build Conflict graph
Acyclic? Topo Sort

Step3

Example serial schedules
Conflict Equiv to S1

| | | | | | | | | | |
|-------|-------|-------|-------|-------|-------|-------|-------------------|-------|-------|
| | T3 | T1 | T4 | T2 | T5 | | SerialSched (SS1) | | |
| w3(C) | w1(A) | w1(B) | r4(B) | w4(E) | r2(A) | r2(A) | w2(D) | r5(D) | w5(E) |
| | T1 | T3 | T2 | T4 | T5 | | SerialSched (SS2) | | |
| w1(A) | w1(B) | w3(C) | r2(A) | r2(A) | w2(D) | r4(B) | w4(E) | r5(D) | w5(E) |

Example with 5 Transactions (2PL)

Schedule S1

| | | | | | | | | | |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| w1(A) | r2(A) | w1(B) | w3(C) | r2(C) | r4(B) | w2(D) | w4(E) | r5(D) | w5(E) |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|

Execute with 2PL

T1 T2 T3 T4 T5

Step 1

X (A)

w1(A)

Step 1.1

Req S(A)

Step 2

X (B)

w1(B)

Unl B, A

Step 3

X (C)

w3(C)

Unl C

Steps 4, 5

Get S(A)

r2(A)

S(C)

r2(C)

Step 6

S(B)

r4(B)

Step 7

X(D)

w2(D)

Unl A, C, D

Step 8

X(E)

w4(E)

Unl B, E

Step 9

S (D)

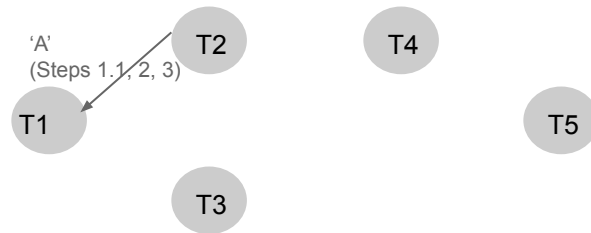
r5(D)

Step 10

X (E)

w5(E)

Unl D, E



Waits- For Graph

Materials to
know

A few
questions

| | |
|------------------------------|---|
| Know your SQL | Select-Project-Aggregate-Having, Nested queries, NULL values Set vs multi-sets |
| How to design 'good' tables? | Find bad FDs, Compute closures/superkeys, Do BCNF decompositions, Find MVDs |
| ACID transactions | Performance #s (disk seeks, memory access) WAL logs for a transaction? How to recover post crash? (Recap: WAL & performance, Why is a log COMMIT fast?) |
| | Is a schedule conflict serializable? How to build Conflict graph? What happens during 2PL execution? (Recap: With a 5 Transaction example) |

(Recap will start with transactions and go backwards)

Example Enrollment table - “v0”

~375
cs145
students

| SID | Class | Room | Time | Lat | Lng |
|---------|--------|------------|------------|------------|-------------|
| 4749732 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 2720942 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 4823984 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 4287594 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2984994 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 8472374 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 4723663 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2478239 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 4763268 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2364532 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2364573 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 3476382 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2347623 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| ... | ... | ... | ... | ... | ... |
| 2364579 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |
| 3476343 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |
| 2322232 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |

~300
cs245
students



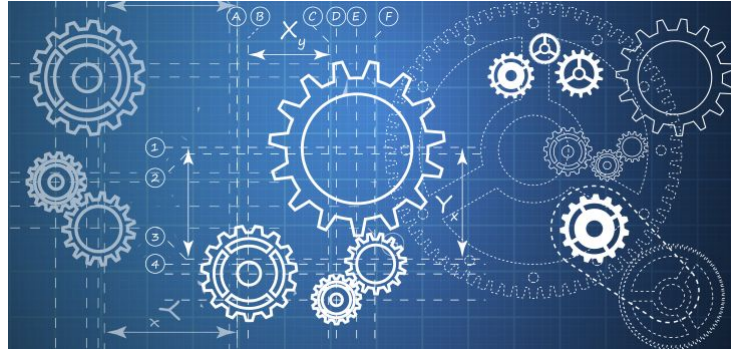
Problems
Repeats?
Room/time change?
Deletes?

Properties
Class -> Room/time
Room -> Lat, Lng

(more compact)

Design Theory

- Design theory is about how to represent your data to avoid *anomalies*.
- Simple algorithms for “best practices”



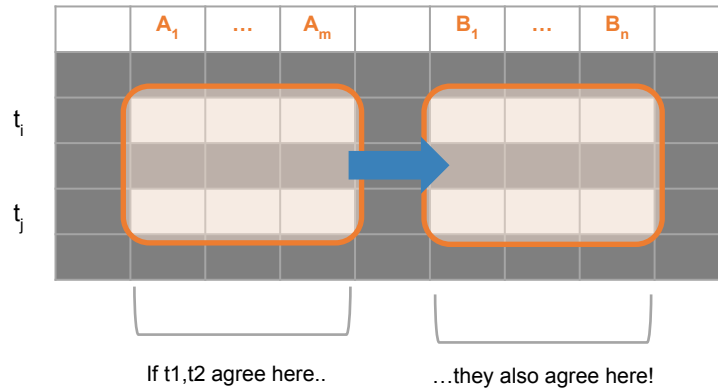
A close-up photograph of a hand holding a blue pen, poised to write on a piece of paper. The hand is wearing a grey, textured sweater. The background is blurred, showing a desk and some papers.

Relational Schema Design

High-level idea

1. Start with some relational *schema*
2. Find out its *functional dependencies (FDs)*
3. Use these to *design a better schema*
One which minimizes the possibility of anomalies

A Picture Of FDs



Defn (again):

Given attribute sets $A = \{A_1, \dots, A_m\}$ and $B = \{B_1, \dots, B_n\}$ in R ,

The **functional dependency** $A \rightarrow B$ on R holds if for **any** t_i, t_j in R :

if $t_i[A_1] = t_j[A_1]$ AND $t_i[A_2] = t_j[A_2]$ AND ... AND $t_i[A_m] = t_j[A_m]$

then $t_i[B_1] = t_j[B_1]$ AND $t_i[B_2] = t_j[B_2]$ AND ... AND $t_i[B_n] = t_j[B_n]$



Finding Functional Dependencies

Equivalent to asking: Given a set of FDs, $F = \{f_1, \dots, f_n\}$, does an FD g hold?

Inference problem: How do we decide?

Answer: Three simple rules called
Armstrong's Rules.

1. **Split/Combine,**
2. **Reduction, and**
3. **Transitivity...** *ideas by picture*

Finding Functional Dependencies

Example:

Products

| Name | Color | Category | Dep | Price |
|--------|-------|----------|--------|-------|
| Gizmo | Green | Gadget | Toys | 49 |
| Widget | Black | Gadget | Toys | 59 |
| Gizmo | Green | Whatsit | Garden | 99 |

Provided FDs:

1. {Name} \rightarrow {Color}
2. {Category} \rightarrow {Department}
3. {Color, Category} \rightarrow {Price}

Which / how many other FDs hold?

Finding Functional Dependencies

Example:

Inferred FDs:

| Inferred FD | Rule used |
|---|-----------------------|
| 4. {Name, Category} → {Name} | Trivial |
| 5. {Name, Category} → {Color} | Transitive (4 → 1) |
| 6. {Name, Category} → {Category} | Trivial |
| 7. {Name, Category} → {Color, Category} | Split/Combine (5 + 6) |
| 8. {Name, Category} → {Price} | Transitive (7 → 3) |

Provided FDs:

1. {Name} → {Color}
2. {Category} → {Dept.}
3. {Color, Category} → {Price}

What's an algorithmic way to do this?

Closure of a set of Attributes

Given a set of attributes A_1, \dots, A_n and a set of FDs F :
Then the closure, $\{A_1, \dots, A_n\}^+$ is the set of attributes B s.t. $\{A_1, \dots, A_n\} \rightarrow B$

Algorithm

Start with $X = \{A_1, \dots, A_n\}$, FDs F .
Repeat until X doesn't change; **do**:
 if $\{B_1, \dots, B_n\} \rightarrow C$ is in F **and**
 $\{B_1, \dots, B_n\} \subseteq X$:
 then add C to X .
Return X as X^+

Example:

$F =$

$\{\text{name}\} \rightarrow \{\text{color}\}$
 $\{\text{category}\} \rightarrow \{\text{department}\}$
 $\{\text{color, category}\} \rightarrow \{\text{price}\}$

Example Closures:

$\{\text{name}\}^+ = \{\text{name, color}\}$
 $\{\text{name, category}\}^+ =$
 $\{\text{name, category, color, dept, price}\}$
 $\{\text{color}\}^+ = \{\text{color}\}$



Keys and Superkeys

A **superkey** is a set of attributes A_1, \dots, A_n s.t.
for *any other* attribute B in R ,
we have $\{A_1, \dots, A_n\} \rightarrow B$

I.e. all attributes are
functionally determined by
a superkey

A **key** is a *minimal* superkey

Meaning that no subset of a
key is also a superkey

Algorithm: For each set of attributes X

1. Compute X^+
2. If $X^+ =$ set of all attributes then
 X is a **superkey**
3. If X is minimal, then it is a **key**

Conceptual Design



For a “mega” table

- Search for “bad” dependencies
- If any, *keep decomposing (lossless) the table into sub-tables* until no more bad dependencies
- When done, the database schema is normalized

Recall: there are several normal forms...



Boyce-Codd Normal Form

BCNF is a simple condition for removing anomalies from relations:

A relation R is **in BCNF** if:

if $\{A_1, \dots, A_n\} \rightarrow B$ is a *non-trivial* FD in R
then $\{A_1, \dots, A_n\}$ is a **superkey** for R

In other words: there are no “bad” FDs



BCNF Decomposition Algorithm

Algorithm: BCNFDecomp(R):

Find a *set of attributes* X s.t.: $X^+ \neq X$ and $X^+ \neq [\text{all attributes}]$

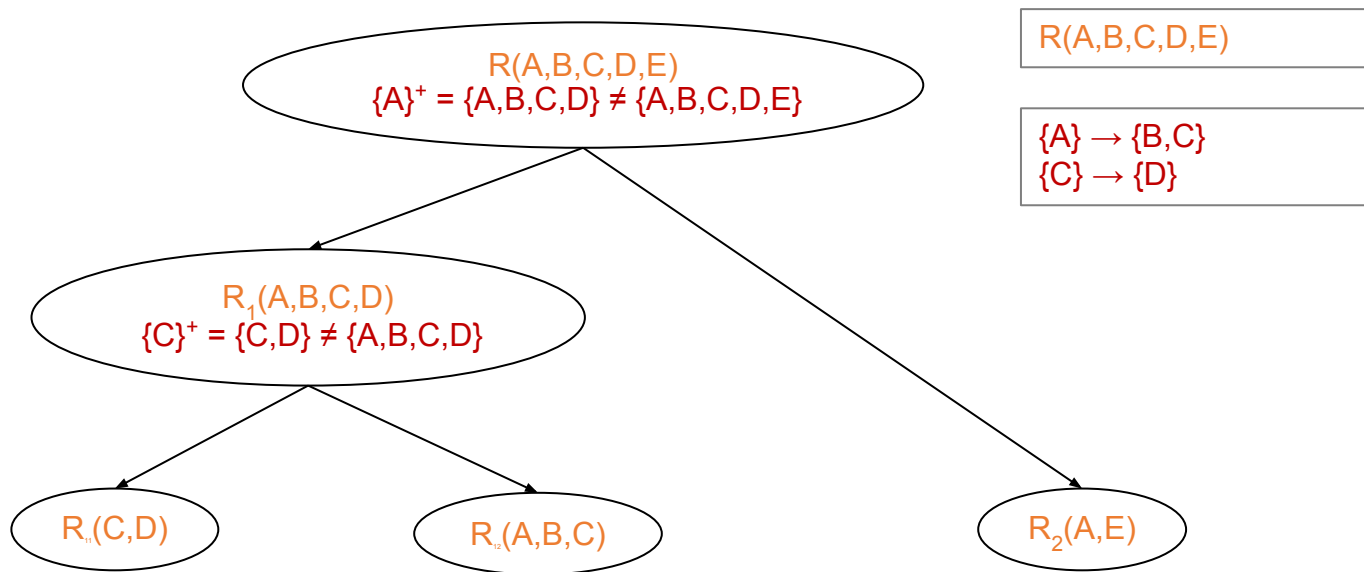
if (not found) then **Return** R

decompose R into $R_1(X^+)$ and $R_2(X \cup \text{Rest})$

Return BCNFDecomp(R_1), BCNFDecomp(R_2)

Proceed recursively until no more “bad” FDs!

Example



Example Enrollment table - “v0”

~375
cs145
students

| SID | Class | Room | Time | Lat | Lng |
|---------|--------|------------|------------|------------|-------------|
| 4749732 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 2720942 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 4823984 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 4287594 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2984994 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 8472374 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 4723663 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2478239 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 4763268 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2364532 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2364573 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 3476382 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2347623 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| ... | ... | ... | ... | ... | ... |
| 2364579 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |
| 3476343 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |
| 2322232 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |

~300
cs245
students



FDs

Class -> Room, Time
Room -> Lat, Lng

(more compact)

Example Enrollment table - “v0”

BCNF decomposition

| SID | Class | Room | Time | Lat | Lng |
|---------|--------|------------|------------|------------|-------------|
| 4749732 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 2720942 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 4823984 | cs 145 | Nvidia Aud | T/R 4:30-6 | 37.4277° N | 122.1742° W |
| 4287594 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2984994 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 8472374 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 4723663 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2478239 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 4763268 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2364532 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2364573 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 3476382 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| 2347623 | cs 145 | Nvidia Aud | T/R 4:30-6 | ... | ... |
| ... | ... | ... | ... | ... | ... |
| 2364579 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |
| 3476343 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |
| 2322232 | cs 245 | Nvidia Aud | T/R 3-4:30 | 37.4277° N | 122.1742° W |

Schema: SID, Class, Room, Time, Lat, Lng

Key

SID, Class

FDs

Class -> Room, Time

Room -> Lat, Lng

BCNF decomposition

1. Find bad FD #1: $\text{Class}^+ \rightarrow \text{Class}, \text{Room}, \text{Time}, \text{Lat}, \text{Lng}$
Decomposed: R1(Class, Room, Time, Lat, Lng) and R2(SID, Class)
2. Find bad FD #2: $\text{Room}^+ \rightarrow \text{Room}, \text{Lat}, \text{Lng}$
Decompose R1 into R11(Room, Lat, Lng) and R12(Class, Room, Time)

⇒ BCNF schema: R2(SID, Class), R12(Class, Room, Time), R11(Room, Lat, Lng)



Example Enrollment table - “v1”

375
cs145
students

| SID | Class |
|---------|--------|
| 4749732 | cs 145 |
| 2720942 | cs 145 |
| 4823984 | cs 145 |
| 4287594 | cs 145 |
| 2984994 | cs 145 |
| 8472374 | cs 145 |
| 4723663 | cs 145 |
| 2478239 | cs 145 |
| 4763268 | cs 145 |
| 2364532 | cs 145 |
| 2364573 | cs 145 |
| 3476382 | cs 145 |
| 2347623 | cs 145 |
| ... | ... |
| 2364579 | cs 245 |
| 3476343 | cs 245 |
| 2322232 | cs 245 |

300
cs245
students

| Class | Room | Time |
|--------|------------|------------|
| cs 145 | Nvidia Aud | T/R 4:30-6 |
| cs 245 | Nvidia Aud | T/R 3-4:30 |
| cs 246 | Nvidia Aud | M/W 3-4:30 |

| Room | Lat | Lng |
|------------|------------|-------------|
| Nvidia Aud | 37.4277° N | 122.1742° W |





“Student 4749732 is taking
Classes = {cs145, cs 245, cs 222},
Hobbies = {Surfing, Music, Astronomy}”

Example:
Student profile
(MVDs)

| SID | Class | Hobby |
|---------|--------|-----------|
| 4749732 | cs 145 | Surfing |
| 4749732 | cs 245 | Surfing |
| 4749732 | cs 222 | Surfing |
| 4749732 | cs 145 | Music |
| 4749732 | cs 245 | Music |
| 4749732 | cs 222 | Music |
| 4749732 | cs 145 | Astronomy |
| 4749732 | cs 245 | Astronomy |
| 4749732 | cs 222 | Astronomy |
| 8472374 | cs 145 | - |
| 8472374 | cs 336 | - |
| ... | | |

Problem

FDs?

Lots of redundancy

Root cause

Conditional independence

given SID: Classes & Hobbies are independent

MVDs: Movie Theatre Example

| | Movie_theater (A) | film_name (B) | Snack (C) |
|-------|-------------------|--|-------------|
| t_1 | Rains 216 | Star Trek: The Wrath of Kahn | Kale Chips |
| t_3 | Rains 216 | Star Trek: The Wrath of Kahn | Burrito |
| | Rains 216 | Lord of the Rings: Concatenated & Extended Edition | Kale Chips |
| t_2 | Rains 216 | Lord of the Rings: Concatenated & Extended Edition | Burrito |
| | Rains 218 | Star Wars: The Boba Fett Prequel | Ramen |
| | Rains 218 | Star Wars: The Boba Fett Prequel | Plain Pasta |

Conditional independence
(between film and snack
given movie theatre)!

$\{A\} \twoheadrightarrow \{B\}$ if for any tuples t_1, t_2 s.t.
 $t_1[A] = t_2[A]$

there is a tuple t_3 s.t.

- $t_3[A] = t_1[A]$
- $t_3[B] = t_1[B]$
- and $t_3[R \setminus B] = t_2[R \setminus B]$

Where $R \setminus B$ is “R minus B” i.e. the
attributes of R not in B.

Example: there is a tuple t_3

$t_3[A] = t_1[A] = \text{'Rains 216'}$

$t_3[B] = t_1[B] = \text{'Star Trek: Wrath of Kahn'}$

$t_3[R \setminus B] = t_2[R \setminus B] = \text{'Burrito'}$

Notes:

1. you can also check the other tuple
2. With more columns, $[R \setminus B]$ will be C + more columns, and $t_3[R \setminus B]$ should = $t_2[R \setminus B]$

Materials to know

A few questions

| | |
|------------------------------|---|
| Know your SQL | Select-Project-Aggregate-Having, Nested queries, NULL values Set vs multi-sets |
| How to design 'good' tables? | Find bad FDs, Compute closures/superkeys, Do BCNF decompositions, Find MVDs |
| ACID transactions | Performance #s (disk seeks, memory access) WAL logs for a transaction? How to recover post crash? (Recap: WAL & performance, Why is a log COMMIT fast?) |
| | Is a schedule conflict serializable? How to build Conflict graph? What happens during 2PL execution? (Recap: With a 5 Transaction example) |

(Recap will start with transactions and go backwards)

General form of Grouping and Aggregation

```
SELECT  S
FROM    R1,...,Rn
WHERE   C1
GROUP BY a1,...,ak
HAVING  C2
```

Evaluation steps:

1. Evaluate **FROM-WHERE**: apply condition C_1 on the attributes in R_1, \dots, R_n
2. **GROUP BY** the attributes a_1, \dots, a_k
3. Apply **HAVING** condition C_2 to each group (may have aggregates)
4. Compute aggregates in **SELECT**, S , and return the result

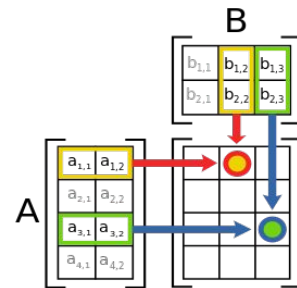


General clarification: Sets vs. Multisets

- In theory, and in any more formal material, by definition all relations are sets of tuples
- In SQL, relations (i.e. tables) are **multisets**, meaning you can have duplicate tuples
 - We need this because intermediate results in SQL don't eliminate duplicates
- If you get confused: just state your assumptions & we'll be forgiving!

Linear Algebra, Declaratively

- Matrix multiplication & other operations = just **joins**!
- The shift from **procedural** to **declarative** programming



```
C = [[0]*p for i in range(n)]
```

```
for i in range(n):  
  for j in range(p):  
    for k in range(m):  
      C[i][j] += A[i][k] * B[k][j]
```



```
SELECT A.i, B.j, SUM(A.x * B.x)  
FROM A, B  
WHERE A.j = B.i  
GROUP BY A.i, B.j;
```

Proceed through a series of instructions

Declare a desired output set



Common SQL Query Paradigms

GROUP BY / HAVING + Aggregators + Nested queries

```
SELECT station_id,  
       COUNT(day) AS nbd  
FROM precipitation,  
     (SELECT day, MAX(precip)  
      FROM precipitation  
      GROUP BY day) AS m  
WHERE day = m.day AND precip = m.precip  
GROUP BY station_id  
HAVING COUNT(day) > 1  
ORDER BY nbd DESC;
```

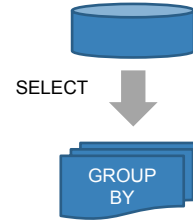
Think about **order***!

**of the semantics, not the actual execution*

Common SQL Query Paradigms

GROUP BY / HAVING + Aggregators + Nested queries

```
SELECT station_id,  
       COUNT(day) AS nbd  
FROM precipitation,  
     (SELECT day, MAX(precip)  
      FROM precipitation  
      GROUP BY day) AS m  
WHERE day = m.day AND precip = m.precip  
GROUP BY station_id  
HAVING COUNT(day) > 1  
ORDER BY nbd DESC;
```

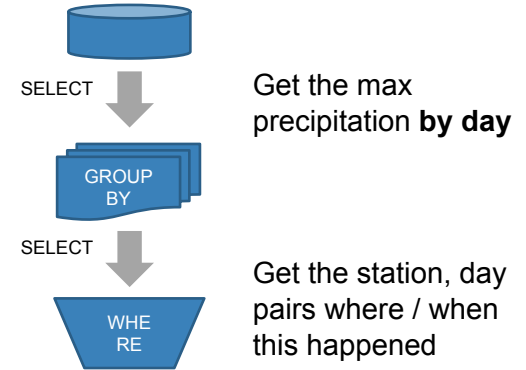


Get the max precipitation **by day**

Common SQL Query Paradigms

GROUP BY / HAVING + Aggregators + Nested queries

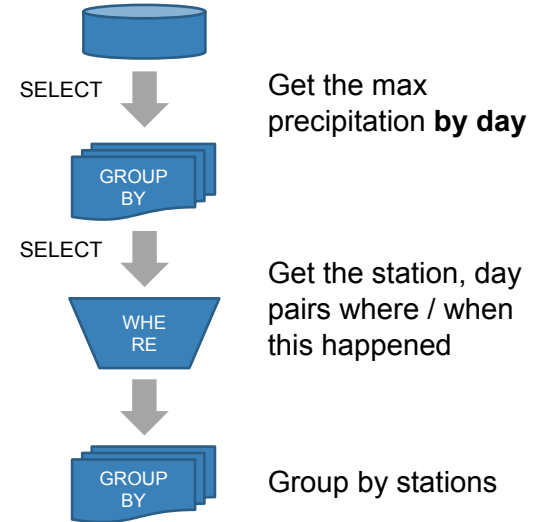
```
SELECT station_id,  
       COUNT(day) AS nbd  
FROM precipitation,  
     (SELECT day, MAX(precip)  
      FROM precipitation  
      GROUP BY day) AS m  
WHERE day = m.day AND precip = m.precip  
GROUP BY station_id  
HAVING COUNT(day) > 1  
ORDER BY nbd DESC;
```



Common SQL Query Paradigms

GROUP BY / HAVING + Aggregators + Nested queries

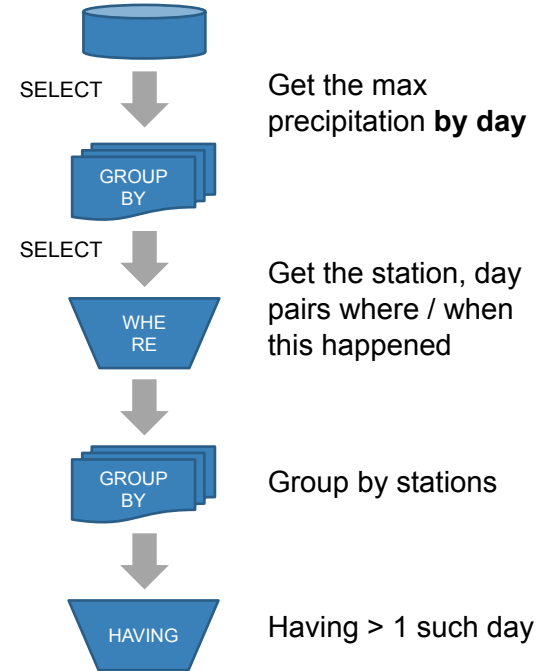
```
SELECT station_id,  
       COUNT(day) AS nbd  
FROM precipitation,  
     (SELECT day, MAX(precip)  
      FROM precipitation  
      GROUP BY day) AS m  
WHERE day = m.day AND precip = m.precip  
GROUP BY station_id  
HAVING COUNT(day) > 1  
ORDER BY nbd DESC;
```



Common SQL Query Paradigms

GROUP BY / HAVING + Aggregators + Nested queries

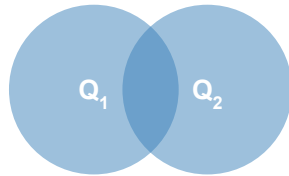
```
SELECT station_id,  
       COUNT(day) AS nbd  
FROM precipitation,  
     (SELECT day, MAX(precip)  
      FROM precipitation  
      GROUP BY day) AS m  
WHERE day = m.day AND precip = m.precip  
GROUP BY station_id  
HAVING COUNT(day) > 1  
ORDER BY nbd DESC;
```





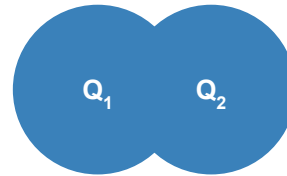
INTERSECT

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
INTERSECT  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```



UNION

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
UNION  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```



EXCEPT

```
SELECT R.A  
FROM R, S  
WHERE R.A=S.A  
EXCEPT  
SELECT R.A  
FROM R, T  
WHERE R.A=T.A
```





Nested queries: Sub-queries Returning Relations

```
Company(name, city)
Product(name, maker)
Purchase(id, product, buyer)
```

```
SELECT c.city
FROM Company c
WHERE c.name IN (
  SELECT pr.maker
  FROM Purchase p, Product pr
  WHERE p.product = pr.name
  AND p.buyer = 'Joe Blow')
```

“Cities where one
can find
companies that
manufacture
products bought
by Joe Blow”

Nested Queries: Operator Semantics

Product(name, price, category, maker)

ALL

```
SELECT name
FROM Product
WHERE price > ALL(
  SELECT price
  FROM Product
  WHERE maker = 'G')
```

Find products that are more expensive than **all products** produced by “G”

ANY

```
SELECT name
FROM Product
WHERE price > ANY(
  SELECT price
  FROM Product
  WHERE maker = 'G')
```

Find products that are more expensive than **any one product** produced by “G”

EXISTS

```
SELECT name
FROM Product p1
WHERE EXISTS (
  SELECT *
  FROM Product p2
  WHERE p2.maker = 'G'
  AND p1.price =
  p2.price)
```

Find products where **there exists some** product with the same price produced by “G”

Nested Queries: Operator Semantics

Product(name, price, category, maker)

ALL

```
SELECT name
FROM Product
WHERE price > ALL(X)
```

Price must be > *all* entries in multiset X

ANY

```
SELECT name
FROM Product
WHERE price > ANY(X)
```

Price must be > *at least one* entry in multiset X

EXISTS

```
SELECT name
FROM Product p1
WHERE EXISTS (X)
```

X must be non-empty

**Note that p1 can be referenced in X (correlated query!)*



Null Values

- *For numerical operations*, NULL -> NULL:
 - If $x = \text{NULL}$ then $4 \cdot (3 - x) / 7$ is still NULL
- *For boolean operations*, in SQL there are three values:

| | | |
|---------|---|-----|
| FALSE | = | 0 |
| UNKNOWN | = | 0.5 |
| TRUE | = | 1 |

- If $x = \text{NULL}$ then $x = \text{"Joe"}$ is UNKNOWN



Null Values

- $C1 \text{ AND } C2 = \min(C1, C2)$
- $C1 \text{ OR } C2 = \max(C1, C2)$
- $\text{NOT } C1 = 1 - C1$

```
SELECT *  
FROM Person  
WHERE (age < 25)  
AND (height > 6 AND weight > 190)
```

Won't return e.g.
(age=20
height=NULL
weight=200)!

Rule in SQL: include only tuples that yield TRUE / 1.0

Null Values

Unexpected behavior:

```
SELECT *  
FROM Person  
WHERE age < 25  
OR age >= 25
```

Some Persons are not included !



```
SELECT *  
FROM Person  
WHERE age < 25  
OR age >= 25  
OR age IS NULL
```

Now it includes all Persons!

Can test for NULL explicitly:

- x IS NULL
- x IS NOT NULL

Materials to
know

A few
questions

| | |
|------------------------------|---|
| Know your SQL | Select-Project-Aggregate-Having, Nested queries, NULL values Set vs multi-sets |
| How to design 'good' tables? | Find bad FDs, Compute closures/superkeys, Do BCNF decompositions, Find MVDs |
| ACID transactions | Performance #s (disk seeks, memory access) WAL logs for a transaction? How to recover post crash? (Recap: WAL & performance, Why is a log COMMIT fast?) |
| | Is a schedule conflict serializable? How to build Conflict graph? What happens during 2PL execution? (Recap: With a 5 Transaction example) |

(Recap will start with transactions and go backwards)