



Lecture 14:

Scale, scale, scale

(Indexing + Sorting, Hashing, Counting)

Announcements

2 guest speakers

DJ Patil, today

Theo Vassilakis, Thursday


Project 2 feedback

- Form: Learning experience? Repetition?
- PDFs for submission
- Motivation: Why colab? Why data visualizations? Why 1 cloud stack?

Metrics Projects 1 and 2 (THANKS!!)


How valuable of a learning experience was this assignment?

200 responses




How difficult was this assignment?

200 responses




How enjoyable was the assignment?

200 responses




How valuable of a learning experience was this assignment?

216 responses




How difficult was this assignment?

216 responses



How enjoyable was the assignment?

216 responses



- Takeaways: P2 got too long and repetitious. PDFs from colab are a pain!
- Tweaks: No PDFs in P3. Split P2 for future.

Project

Goal #1

Full data cycle
on cloud

(recap from
Lecture 1)

1 Run queries on public datasets

2 Explore/Visualize public datasets

3 Predict using Machine Learning

 Google Cloud Platform



Data Exploration
Colab



BigQuery (SQL)



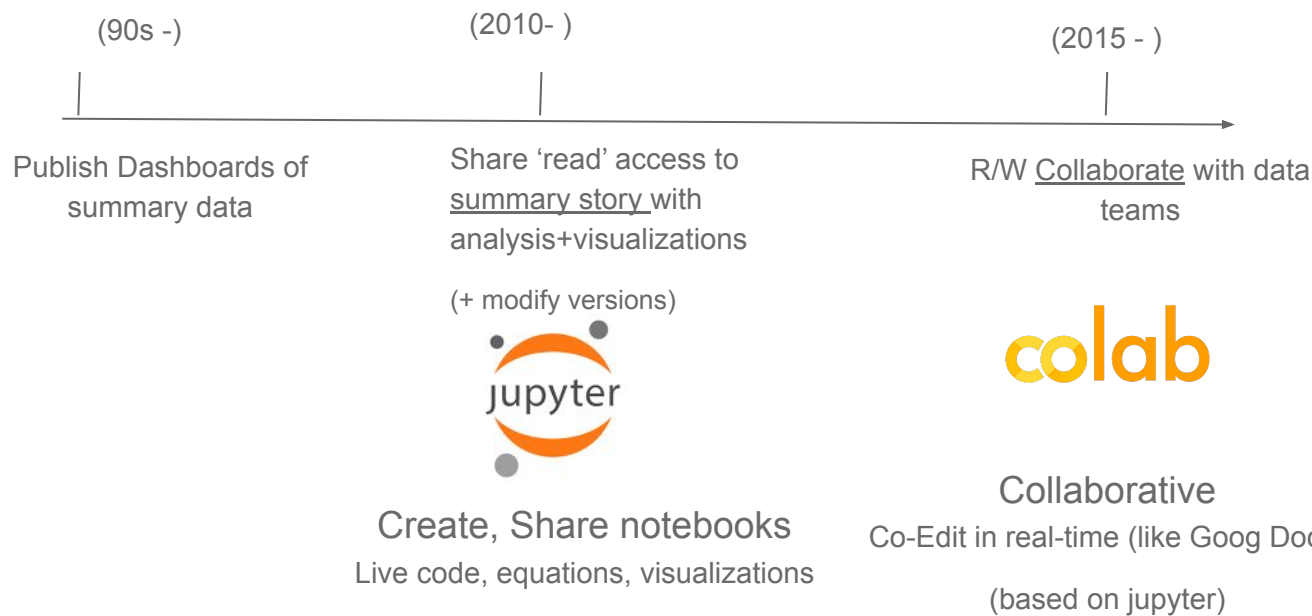
Cloud Machine
Learning

Notes

- Concepts are the key. Easy to go 'cloud hopping'
- Key goal: Do (a) deep data cycle, (b) for 10s-100s of GBs of real data

Goal #2

Sharing data
cycle from 1 to
teams of 10s to
1000s



Summary

When you finish P3...

- Full Query-Visualize-Learn data cycle on cloud stack
 - Experience working on 10s-100s of GBs of data on machine cluster
 - Gaining an ability to grok multiple real datasets (ncaa, github, wbank, ...)
 - Working with popular SQL/Visualization/Learning libraries
 - Gaining intuition of scale, performance, pain and costs in running 100s of queries
- Collaborate in small teams for r/w access to data analysis (e.g., colab)
- Hop-Skip-Jump to any cloud stack or application based on same concepts

⇒ Strong suggestion

Do a great job on P3. Link off your CV/resume.

This week

Scale, Scale, Scale

How to read/write indices?

Sorting, Counting, Hashing
(for RAM, Disk, Clusters)

B+ Tree Index Search [recap]

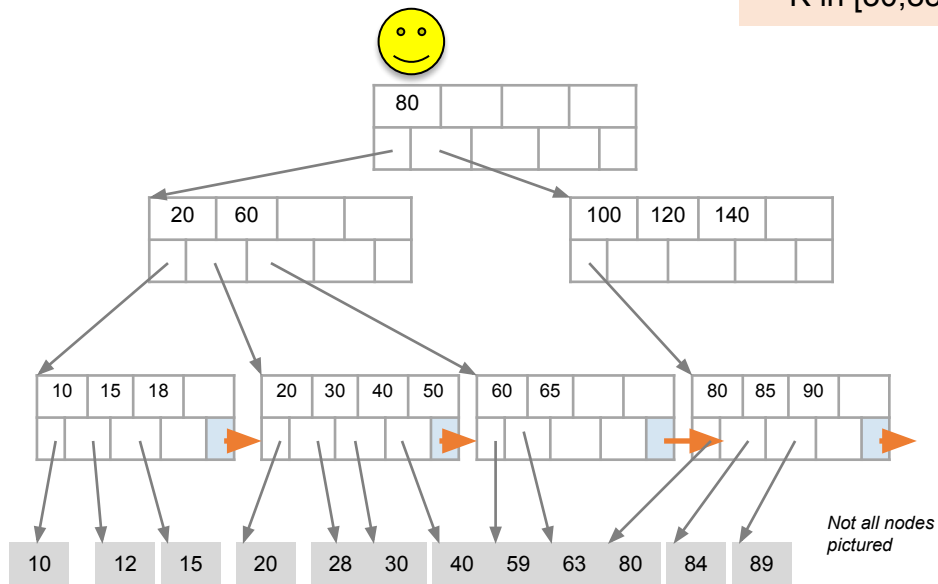
30 < 80

30 in [20,60)

30 in [30,40)

To the data!

K in [30,85)?




Simple Cost Model for Search


- Let:
 - f = fanout, which is in $[d+1, 2d+1]$ (*we'll assume it's constant for our cost model...*)
 - N = the total number of *pages* we need to index
 - F = fill-factor (usually $\approx 2/3$)
- Our B+ Tree needs to have room to index N / F pages!
 - We have the fill factor in order to leave some open slots for faster insertions
- What height (h) does our B+ Tree need to be?
 - $h=1 \rightarrow$ Just the root node- room to index f pages
 - $h=2 \rightarrow f$ leaf nodes- room to index f^2 pages
 - $h=3 \rightarrow f^2$ leaf nodes- room to index f^3 pages
 - ...
 - $h \rightarrow f^{h-1}$ leaf nodes- room to index f^h pages!

\rightarrow We need a B+ Tree of height $h = \left\lceil \log_f \frac{N}{F} \right\rceil$!

Search cost of B+ Tree (on RAM + Disk)




Read 1st levels
Into RAM buffer



$$1 + f + f^2 + f^3 + \dots \leq B$$

Keep 1st L_B levels in
RAM of size B

Rest of index on
disk



Algorithm: B+ Search
 - Read 1 page per level
 - Pages in RAM are free
 - Read 1 page for record

$$\text{IO Cost: } \left\lceil \log_f \frac{N}{f} \right\rceil - L_B + 1$$

$$\text{where } B \geq \sum_{l=0}^{L_B-1} f^l$$



Simple Cost Model for Search

- To do range search, we just follow the horizontal pointers
- The IO cost is that of loading additional leaf nodes we need to access + the IO cost of loading each **page** of the results- we phrase this as “Cost(OUT)”

$$\text{IO Cost: } \left\lceil \log_f \frac{N}{F} \right\rceil - L_B + \text{Cost}(\text{OUT})$$

$$\text{where } B \geq \sum_{l=0}^{L_B-1} f^l$$



Fast Insertions & Self-Balancing

- We won't go into specifics of B+ Tree insertion algorithm, but has several attractive qualities:
 - ~ Same cost as exact search
 - **Self-balancing:** B+ Tree remains **balanced** (with respect to height) even after insert

B+ Trees also (relatively) fast for single insertions!
However, bottleneck if many insertions (if fill-factor slack is used up...)


A close-up photograph of a person's hand holding a blue pen, poised to write on a white sheet of paper. The hand is wearing a grey, textured sweater. The background is blurred, showing more of the paper and the pen.

Clustered Indexes

An index is **clustered** if the underlying data is ordered in the same way as the index's data entries.

Clustered vs. Unclustered Index


E.g. Person(name, age, SSN)



Clustered

(e.g., Sorted by SSN,
Index on SSN)

Index Entries



Data Records

Unclustered

(e.g. index for age)




Clustered vs. Unclustered Index

- Recall that for a disk with block access, **sequential IO is much faster than random IO**
- For exact search, no difference between clustered / unclustered
- For range search over R values: difference between **1 random IO + R sequential IO**, and **R random IO**:
 - A random IO costs ~ 10ms (sequential much much faster)
 - For R = 100,000 records- **difference between ~10ms and ~17min!**

Big Scaling (with Indexes)

Roadmap



Primary data structures/algorithms

Hashing

HashTables
($\text{hash}_i(x)$)

Sorting

BucketSort, QuickSort
MergeSort

Counting

HashTable + Counter
($\text{hash}_i(\text{key}) \rightarrow \langle \text{count} \rangle$)

MergeSortedFiles
SortFiles

?????

MergeSortedFiles
SortFiles




Recall: Hashing

- **Magic of hashing:**
 - A hash function h_B maps into $[0, B-1]$
 - And maps nearly uniformly
- A hash **collision** is when $x \neq y$ but $h_B(x) = h_B(y)$
 - Note however that it will never occur that $x = y$ but $h_B(x) \neq h_B(y)$
- E.g., hash on any attribute A

Hashing for scale

- $h_i(x)$ % one of below
 - B = memory buf (h_B)
 - M = machines, shards (h_M)
 - $\langle B, M \rangle$ = memory buf, machines ($h_{B, M}$)

Algorithm: Where is record 'x' for read/write
if $\text{Buf}[h_i(x)]$ matches 'x' // check if really x?
Use record



- Multiple hash functions (uncorrelated to spread data)
 - $h_i(x), h_{i+1}(x), h_{i+2}(x), h_{i+3}(x), \dots$


Big Scaling (with Indexes)

Roadmap

Primary data structures/algorithms

Hashing

HashTables
($\text{hash}_i(x)$)



Hashes for disk location
($\text{hash}_i(x)$)



Hashes for machines,
shards
($\text{hash}_i(x)$)

Sorting

BucketSort, QuickSort
MergeSort

MergeSortedFiles
SortFiles

MergeSortedFiles
SortFiles

Counting

HashTable + Counter
($\text{hash}_i(\text{key}) \rightarrow \langle \text{count} \rangle$)

?????

Summary

Scale with Sorting, Hashing, Counting

Sorting

- Indices over sorted files (B+ trees) --
lookups/ranges
- Index pages with big fanout


Hashing

- Hash Indices over data records

Clustered vs non-clustered indices

Counting?

Counting product views for billion products



Nespresso Vertuo Coffee and Espresso Machine Bundle with Aeroccino Milk Frother by Breville, Red

by Breville
★★★★★ 980 customer reviews
259 answered questions
Amazon's Choice for "nespresso machine red"

List Price: \$249.95
Price: **\$189.96** ✓prime | FREE One-Day
You Save: \$59.99 (24%)
Your cost could be \$179.96. Eligible customers get a \$10 bonus when reloading \$100.

Free Amazon product support included

Style Name: **Nespresso by Breville**

Nespresso Nespresso by Breville

Color: Red



Counting popular product-pairs

Customers who viewed this item also viewed these products



Dualit Food XL1500
Processor
\$560

Add to cart



Kenwood kMix Manual
Espresso Machine
★★★★★
\$250

Select options



Weber One Touch Gold
Premium Charcoal
Grill-57cm
\$225

Add to cart



NoMU Salt Pepper and
Spice Grinders
\$3

View options

Problem: Number of counters = $O(B*B)$
(We'll discuss next class)

Hints

- most counts will be sparse
- What solves 'all' known CS problems?